

Lab 1 Report

Video

<https://youtu.be/B8JlzyTJXQI>

Changes

The following is a list of files that have been changed, along with what was changed in each file.

Makefile

```
132     $U/_grind\  
133     $U/_wc\  
134     $U/_zombie\  
135 +   $U/_lab1_test\  
136  
137     fs.img: mkfs/mkfs README $(UPROGS)  
138           mkfs/mkfs fs.img README $(UPROGS)
```

The only thing that was changed in this file was the addition of the lab1_test user program.

kernel/defs.h

```
8  struct sleeplock;
9  struct stat;
10 struct superblock;
11 + struct pinfo;
12
13 // bio.c
14 void      binit(void);
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64 void*      kalloc(void);
65 void      kfree(void *);
66 void      kinit(void);
67 + int      get_free(void);
68
69 // log.c
70 void      initlog(int, struct superblock*);
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108 int      either_copyout(int user_dst, uint64 dst, void *src, uint64 len);
109 int      either_copyin(void *dst, int user_src, uint64 src, uint64 len);
110 void      procdump(void);
111 + void      add_call(int);
112 + int      sysinfo(int);
113 + int      procinfo(struct pinfo*);
114
115 // swtch.S
116 void      swtch(struct context*, struct context*);
```

New definitions of functions that were created throughout the lab were added to this file. Specifically: pinfo - for the procinfo syscall, get_free - a helper function to get the number of free pages, add_call - another helper function to count the number of syscalls, sysinfo - the first syscall we were to implement, and procinfo - the second syscall we implemented.

kernel/kalloc.c

```
80      memset((char*)r, 5, PGSIZE); // fill with junk
81      return (void*)r;
82  }
83  +
84  + int get_free(void) {
85  +     struct run *r = kmem.freelist;
86  +     if (r == 0) return 0; // Return if null
87  +
88  +     int num = 0;
89  +
90  +     acquire(&kmem.lock);
91  +     while (r != 0) {
92  +         num++;
93  +         r = r->next;
94  +     }
95  +     release(&kmem.lock);
96  +
97  +     return num;
98  + }
```

This is where the get_free() helper function was implemented that counts the number of free memory pages. It starts with the address of `kmem.freelist` and iterated over it until it

reaches a null pointer. It takes the kmem lock as well to avoid any possible conflicts during execution.

kernel/proc.c

```
55     initlock(&p->lock, "proc");
56     p->state = UNUSED;
57     p->kstack = KSTACK((int) (p - proc));
58 +   p->syscalls = 0;
59 }
60 }
61

682     printf("\n");
683 }
684 }

685 +
686 + static int num_calls = 0;
687 +
688 + void add_call(int num) {
689 +     num_calls++;
690 + }
691 +
692 + int sysinfo(int in) {
693 +     if (in == 0) {
694 +         int count = 0;
695 +         struct proc *p;
696 +
697 +         for(p = proc; p < &proc[NPROC]; p++) {
698 +             if (proc->state != UNUSED) count++;
699 +         }
700 +         return count;
701 +     } else if (in == 1) {
702 +         return num_calls;
703 +     } else if (in == 2) {
704 +         return get_free();
705 +     }
706 +     return -1;
707 + }

709 + int procinfo(struct pinfo *in) {
710 +     if (in == 0) return -1;
711 +
712 +     struct proc *p = myproc();
713 +
714 +     if (p == 0) return -1;
715 +
716 +     struct pinfo data;
717 +
718 +     int pg_count = p->sz >> PGSHIFT; // pg_count = p->sz / PGSIZE
719 +     if (p->sz % PGSIZE != 0) pg_count++; // Account for loss of decimal
720 +
721 +     data.ppid = p->pid;
722 +     data.page_usage = pg_count;
723 +     data.syscall_count = p->syscalls;
724 +
725 +     return copyout(p->pagetable, (uint64)in, (char *)&data, sizeof(struct pinfo));
726 + }
```

This is where the syscalls, sysinfo and procinfo, were implemented. Additionally the helper function add_call was implemented here. The number of syscalls that a process has made is also initialized to 0.

For the **sysinfo** call, and depending on the value of the argument passed to it, it will either: a) Count the used processes, b) return the number of syscalls that have been made by using the **num_calls** variable that is incremented whenever a syscall is made, or c) return the number of free memory pages by using the **get_free** helper function.

kernel/proc.h

```
104     struct file *ofile[NOFILE]; // Open files
105     struct inode *cwd;           // Current directory
106     char name[16];               // Process name (debugging)
107 +   uint64 syscalls;
108 + };
109 +
110 + struct pinfo {
111 +   int ppid;
112 +   int syscall_count;
113 +   int page_usage;
114 };
```

In this header file, the number of syscalls was added to the PCB struct, and the pinfo struct was defined. The `syscalls` element is added and used to keep track of the number of syscalls that a process has made.

kernel/syscall.c

```
101     extern uint64 sys_link(void);
102     extern uint64 sys_mkdir(void);
103     extern uint64 sys_close(void);
104 + extern uint64 sys_sysinfo(void);
105 + extern uint64 sys_procinfo(void);
106
107     // An array mapping syscall numbers from syscall.h
108     // to the function that handles the system call.
109
110
111     [SYS_link]    sys_link,
112     [SYS_mkdir]   sys_mkdir,
113     [SYS_close]   sys_close,
114 + [SYS_sysinfo]  sys_sysinfo,
115 + [SYS_procinfo] sys_procinfo,
116 };
117
118 void
119
120
121     // Use num to lookup the system call function for num, call it,
122     // and store its return value in p->trapframe->a0
123     p->trapframe->a0 = syscalls[num]();
124
125 +
126 + // Increment syscall count AFTER doing the syscall
127 + // to avoid counting the current call.
128 + add_call(num);
129 + p->syscalls++;
130
131 } else {
132     printf("%d %s: unknown sys call %d\n",
133           p->pid, p->name, num);
134 }
```

In this file, the accounting of syscalls was added (to the `syscall` function) as well as the 2 new syscalls added to the array of syscalls. The `add_call` function will increment the total number of syscalls that have been made, and `p->syscalls++` uses the `syscalls` element of the process to keep track of syscalls on a per-process basis.

kernel/syscall.h

```
20  #define SYS_link 19
21  #define SYS_mkdir 20
22  #define SYS_close 21
23  + #define SYS_sysinfo 22
24  + #define SYS_procinfol 23
```

In this file we added 2 new entry numbers for the 2 new syscalls that we created.

kernel/sysproc.c

```
89  release(&tickslock);
90  return xticks;
91  }
92  +
93  + uint64 sys_sysinfo(void) {
94  +  int in;
95  +
96  +  argint(0, &in);
97  +  return sysinfo(in);
98  + }
99  +
100 + uint64 sys_procinfol(void) {
101 +  uint64 in;
102 +  argaddr(0, &in);
103 +  procinfo((void *)in);
104 +  return 0;
105 + }
```

In this file we added the “glue” functions that process the arguments that collect the call arguments and call the syscall function themselves.

user/user.h

```
1  struct stat;  
2  + struct pinfo;  
3  
4  // system calls  
5  int fork(void);  
  
23 char* sbrk(int);  
24 int sleep(int);  
25 int uptime(void);  
26 + int sysinfo(int);  
27 + int procinfo(struct pinfo*);  
28  
29 // ulib.c  
30 int stat(const char*, struct stat*);
```

In this file the 2 new syscalls so they are usable in user-mode programs.

user/usys.pl

```
36 entry("sbrk");  
37 entry("sleep");  
38 entry("uptime");  
39 + entry("sysinfo");  
40 + entry("procinfo");
```

In this file 2 new lines were added to account for the 2 new syscalls that were added.

Note

We *technically* omitted the `lab1_test` file. Because we did not modify this file, and because it is the same for every group, we did not find it important enough to include.

Syscall Processes

User-Mode to Kernel

We first start at the user-level with a call to the `sysinfo` syscall. Using the definitions in `user/usys.pl` and `user/user.h`, the correct syscall index is found (via `kernel/syscall.h`) This index is then used in the `syscall` function (defined in

`kernel/syscall.c`) to index into the `syscalls` array (a “mapping” of syscall numbers to their respective functions). Before we run the system call itself, we first need to collect the arguments to pass to it. This is done in their respective `sys_` functions; `sys_sysinfo` and `sys_procinfo` for the 2 syscalls that we implemented. After the correct arguments have been collected, they are passed to the syscalls themselves. In the next 2 sections we discuss the specific implementations of the system calls that we implemented that happen before the return to user-mode.

sysinfo

For the `sysinfo` we have 4 different return values depending on the incoming argument. If `in` is `0` we count the number of active processes. We consider a running process to be one that does not have the state of `UNUSED` in the list of process blocks. If `in` is `1`, then we return the value of `num_calls`. `num_calls` is incremented every time that a syscall is made. If `in` is `2`, then we invoke the `get_free` helper function. In this function we start with the address of `kmem.freelist` and iterate through it until we reach `NULL` to signify the end of the list. At the end we return the number of free pages that were found which is then returned by `sysinfo`. Finally, if `in` is none of the above values, then we return `-1`.

procinfo

For `procinfo` we start with getting the current process’ PCB and creating a new `pinfo` struct. We calculate the number of pages being used by getting the size of the processes memory through the `sz` element of the PCB. Then then divide that by the `PGSIZE` (which we accomplish by doing a right shift `PGSHIFT` number of times), and account for non-whole pages being used by checking `p->sz % PGSIZE != 0`. If this is true then we increment the page count. With the number of pages being used, we can populate the rest of the `pinfo` struct with the process ID (from `p->pid`), the `page_usage` (using the page count calculated earlier), and `syscall_count` by using the `p->syscalls` entry of the PCB. Finally we use the `copyout` function to copy the “kernel pinfo struct” to the user mode address from the pointer argument.

Kernel to User-Mode

After the syscall returns, we increment the number of syscalls made on a per-process and global basis (in the `syscall` function in `kernel/syscall.c`). After that function returns, the value is propagated to the user mode and execution of the user program continues.

Contributions

For this lab, as we formed a group mid-assignment, we each did our own implementation and compared results at the end. Additionally we also both contributed to this report.